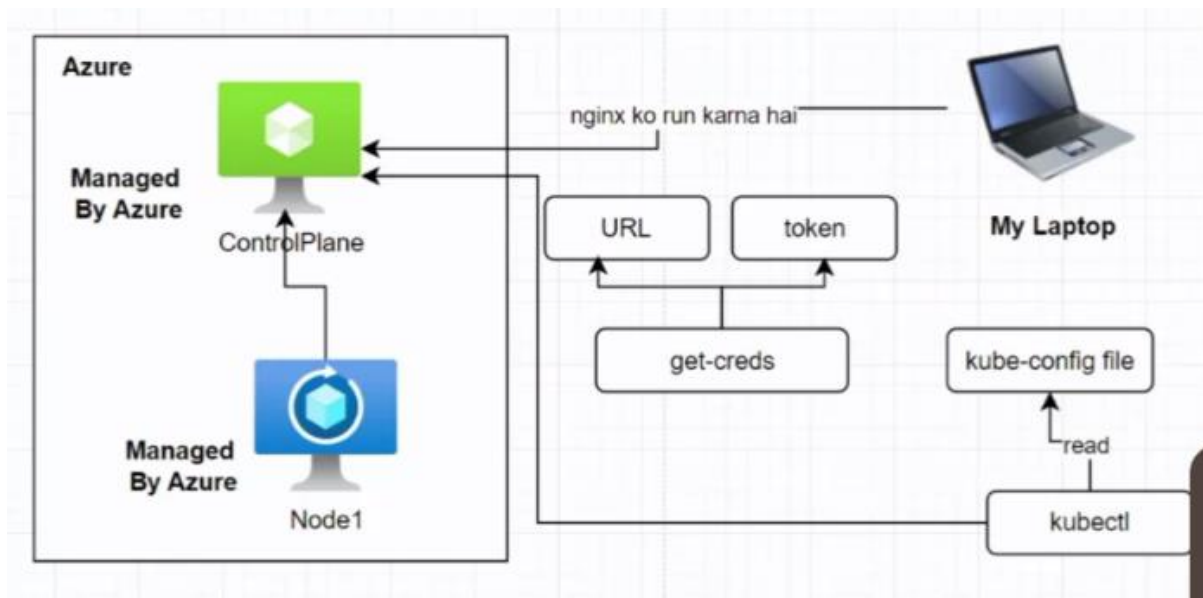


13 October

- 1) Kubernetes is a paas service
- 2) Credentials of cluster api address have been stored in kubectl file
- 3) On my laptop 2 things are stored – kubelet and kubeproxy
- 4)



5) **kubectl run --help**

```
PS C:\Users\HP> kubectl run --help
Create and run a particular image in a pod.

Examples:
# Start a nginx pod
kubectl run nginx --image=nginx

# Start a hazelcast pod and let the container expose port 5701
kubectl run hazelcast --image=hazelcast/hazelcast --port=5701

# Start a hazelcast pod and set environment variables "DNS_DOMAIN=cluster" and "POD_NAMESPACE=default" in the
container
kubectl run hazelcast --image=hazelcast/hazelcast --env="DNS_DOMAIN=cluster" --env="POD_NAMESPACE=default"
```

6) **kubectl run nginx --image=nginx**

```
Use "kubectl options" for a list of global command-line options (applies to all commands).
PS C:\Users\HP> kubectl run nginx --image=nginx
pod/nginx created
PS C:\Users\HP>
```

Now firstly kubectl gone to kube-config file and read the api server url and token and kubectl send the information about installing of nginx on api server (in control plane). Then api server go to scheduler and ask which node is free then after getting free node the kubelet in node will be asked to run nginx pod.

7) Now we have to check what is made in our cluster then how we will see

```
explain get documentation for a resource
get Display one or many resources
```

8) **kubectl get pods**

```
See 'kubectl get -h' for help and examples
PS C:\Users\HP> kubectl get --help
Display one or many resources.

Prints a table of the most important information about
selector and the --selector flag. If the desired resource
namespace if you don't specify any namespace.

By specifying the output as 'template' and providing a
the attributes of the fetched resources.

Use "kubectl api-resources" for a complete list of supported
resources.


Examples:
# List all pods in ps output format
kubectl get pods
```

```
PS C:\Users\HP> kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           9m31s
PS C:\Users\HP>
```

workload = A workload is an application running on Kubernetes.

Pods are the smallest deployable units of computing that you can create and manage in Kubernetes.

Pod = group of one or more containers, with shared storage and network resources



Pause container will store the Pod IP and pod namespace

IP is assigned to POD. Using that we will access containers.

9) Inside pod all containers will use shared networking concept as they will use same ip address and network. And will also share the same storage also

**ek container dusre container
ko localhost se access kar
paega direct...**

Shared Network

Shared Storage

Pod

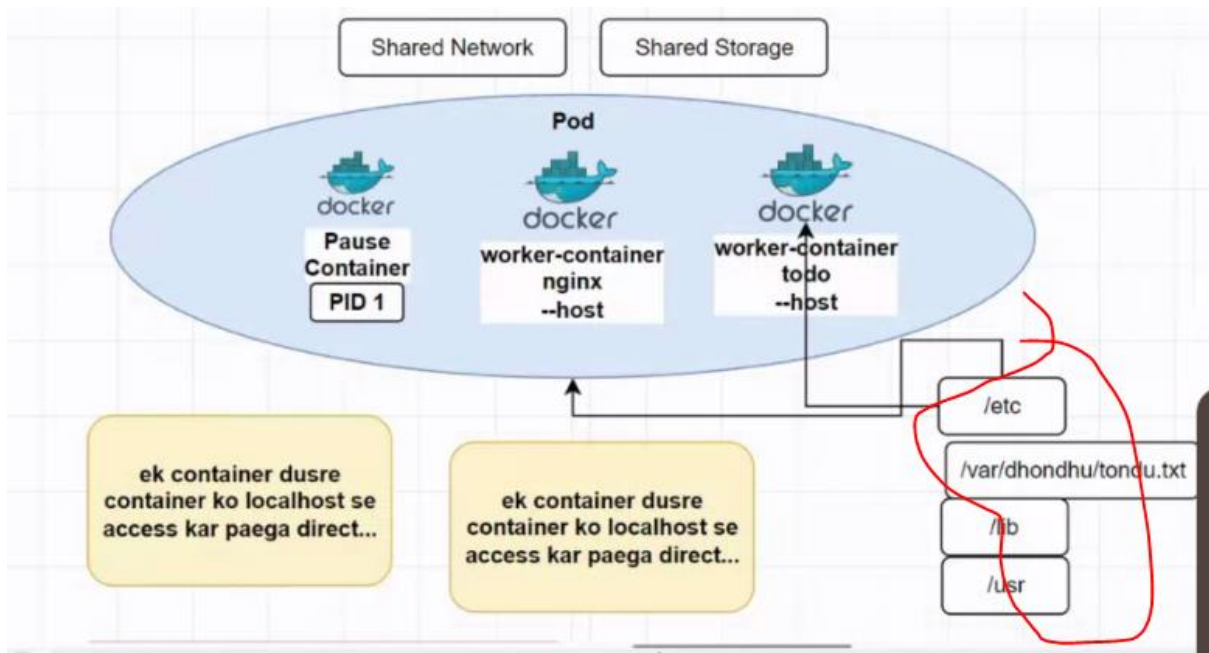


docker
Container 1
--host

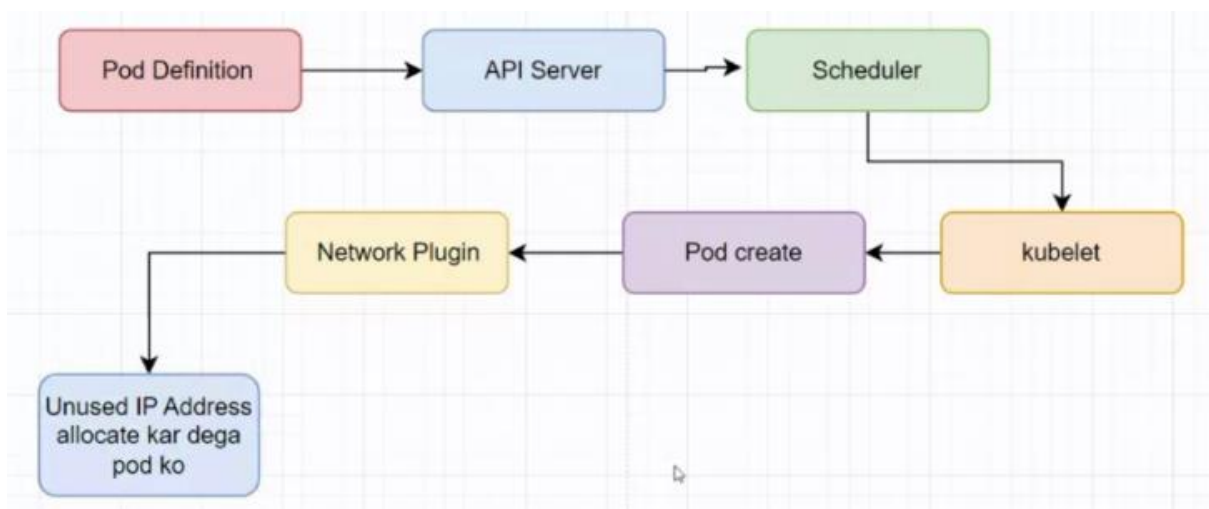
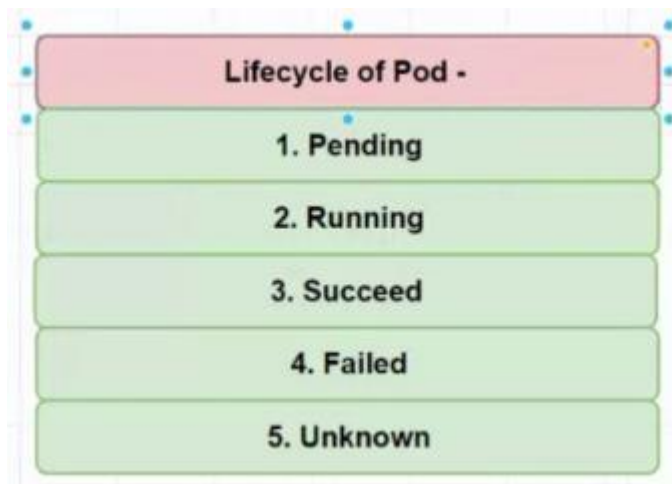


docker
Container 2
--host

HOST - IP

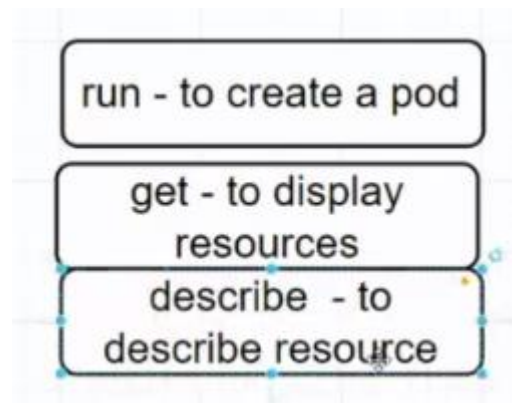


10) pod lifecycle



*) How ip is assigned to plugin = by cni plugin (Container Network Interface (CNI))

11)



```
kubectl describe TYPE NAME_PREFIX
```



```
kubectl describe pod nginx
```

12) **kubectl describe pod nginx**

```

PS C:\Users\HP> kubectl describe pod nginx
Name: nginx
Namespace: default
Priority: 0
Service Account: default
Node: aks-nodepool13-10856213-vmss000000/10.224.0.4
Start Time: Wed, 16 Oct 2024 16:32:57 +0530
Labels: run=nginx
Annotations: <none>
Status: Running
IP: 10.244.1.68
IPs:
  IP: 10.244.1.68
Containers:
  nginx:
    Container ID: containerd://fb1dbb1c654a0e63ed89b26471ae9ef650901d6f222c91061fc20b4a0504e097
    Image: nginx
    Image ID: docker.io/library/nginx@sha256:d2eb56950b84efe34f966a2b92efb1a1a2ea53e7e93b94cdf45a27cf3cd47fc0
    Port: <none>
    Host Port: <none>
    State: Running
      Started: Wed, 16 Oct 2024 16:33:03 +0530
    Ready: True
    Restart Count: 0
    Environment: <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-lsbpv (ro)
Conditions:
  Type              Status
  PodReadyToStartContainers  True
  Initialized          True
  Ready                True
  ContainersReady       True
  PodScheduled          True
Volumes:
  kube-api-access-lsbpv:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events: <none>

```

13) Now since we need to run nginx so we will do port forwarding

```

# Listen on port 8888 locally, forwarding to 5000 in the pod
kubectl port-forward pod/mypod 8888:5000

```

14) **kubectl port-forward pod/nginx 9999:80**

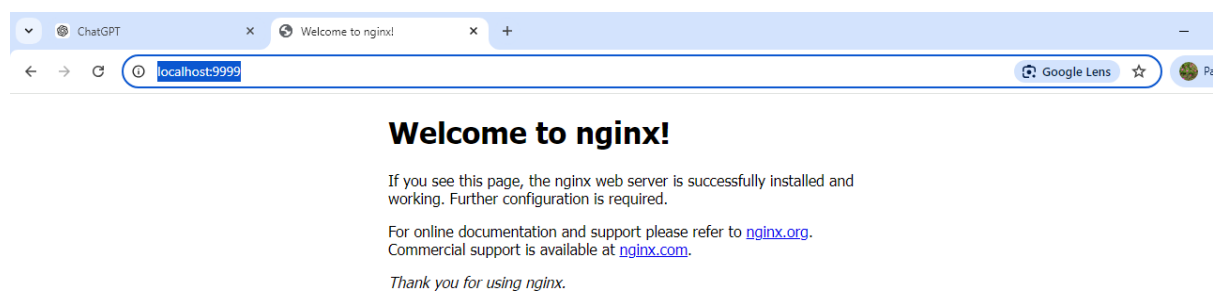
```

PS C:\Users\HP> kubectl port-forward pod/nginx 9999:80
Forwarding from 127.0.0.1:9999 -> 80
Forwarding from [::1]:9999 -> 80

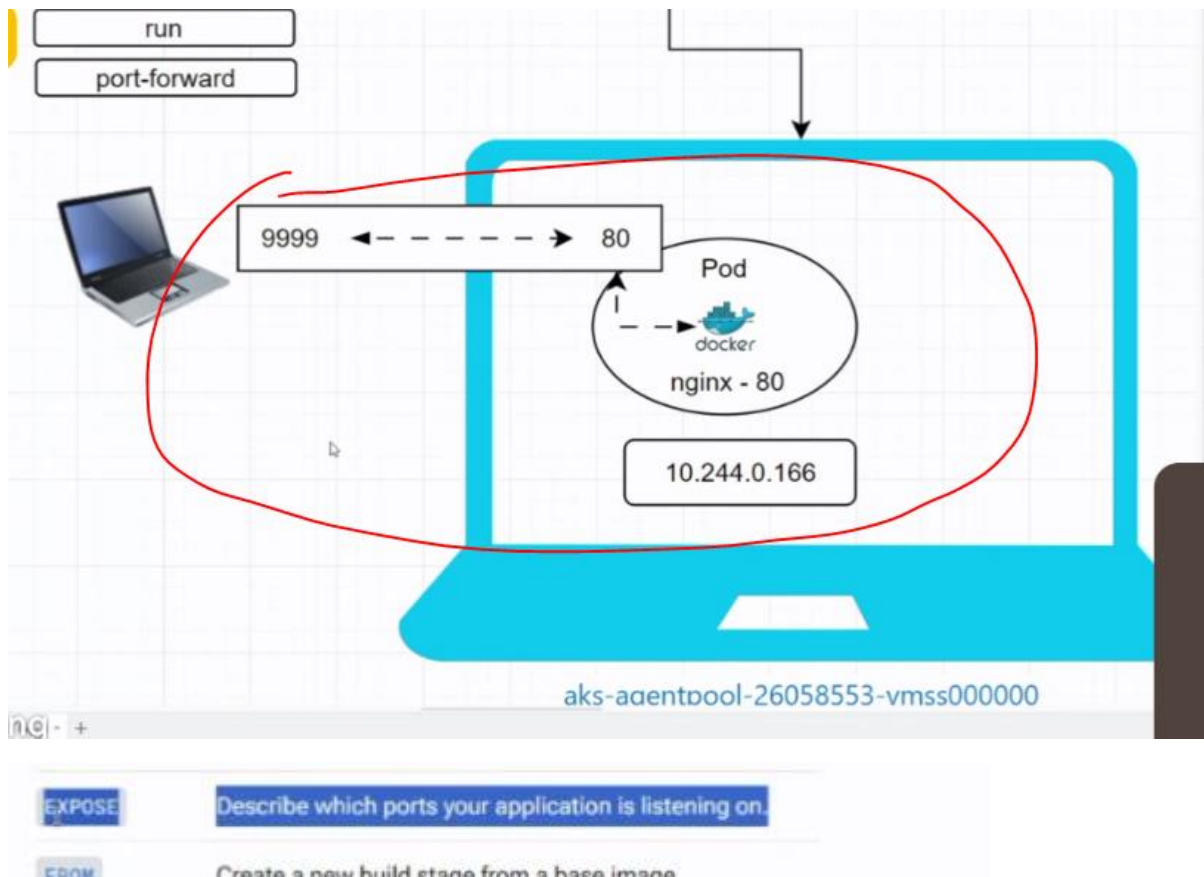
```

15) Run in browser – localhost:9999

<http://localhost:9999/>



16) Since tunnel is created from host to pod and then pod to container. So in pod port syncing between pod and container is already done by **expose command** in docker file. Or the tunnel between pod and container is already made by expose command in docker file.



17) **kubect1 get nodes** – how many nodes are running

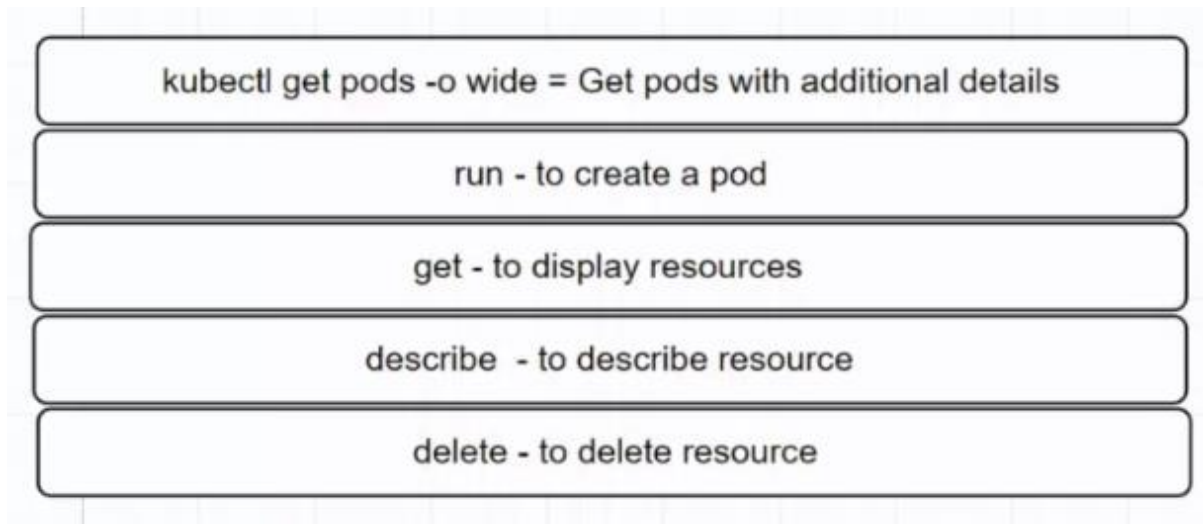
```
PS C:\Users\HP> kubect1 get nodes
NAME                                STATUS    ROLES    AGE    VERSION
aks-agentpool-95341565-vmss000000  Ready    <none>    28m    v1.29.8
aks-nodepool14-95341565-vmss000000  Ready    <none>    28m    v1.29.8
PS C:\Users\HP>
```

18) **kubect1 get pods -o wide** = more info about pod


```
PS C:\Users\HP> kubect1 get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE    IP          NODE                                NOMINATED NODE   READINESS GATES
nginx   1/1     Running   0           18m    10.244.1.240  aks-agentpool-95341565-vmss000000  <none>           <none>
```

19) **kubect1 delete pod nginx** - to delete the pod

```
PS C:\Users\HP> kubect1 delete pod nginx
pod "nginx" deleted
PS C:\Users\HP>
```

20) Why workload run on worker nodes not on master computer? Taint laga hota hai



+++++

AGENDA – CREATE A NAMESPACE

1) **kubectl get ns** – namespace

```
PS C:\Users\HP> kubectl get ns
NAME                STATUS    AGE
default              Active    49m
kube-node-lease      Active    49m
kube-public          Active    49m
kube-system          Active    49m
PS C:\Users\HP>
```

2) **kubectl get pods -n kube-system**


```
PS C:\Users\HP> kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
azure-cns-clmfc	1/1	Running	0	52m
azure-cns-sk59x	1/1	Running	0	51m
azure-ip-masq-agent-fnkcfc	1/1	Running	0	51m
azure-ip-masq-agent-mjsssf	1/1	Running	0	52m
cloud-node-manager-9fp78	1/1	Running	0	52m
cloud-node-manager-sxspl	1/1	Running	0	51m
coredns-597bb9d4db-8c6fc	1/1	Running	0	52m
coredns-597bb9d4db-q6fbp	1/1	Running	0	51m
coredns-autoscaler-689db4649c-1czt9	1/1	Running	0	52m
csi-azuredisk-node-768g5	3/3	Running	0	51m
csi-azuredisk-node-lgb4j	3/3	Running	0	52m
csi-azurefile-node-jspqv	3/3	Running	0	51m
csi-azurefile-node-pvcfb	3/3	Running	0	52m
eraser-controller-manager-6d5c64f9c8-2vrc7	1/1	Running	0	50m
konnektivity-agent-7b8fd8867d-nzvbq	1/1	Running	0	34m
konnektivity-agent-7b8fd8867d-zgfqx	1/1	Running	0	34m
kube-proxy-92gfg	1/1	Running	0	52m
kube-proxy-qx54j	1/1	Running	0	51m
metrics-server-f46f56d7b-9clr4	2/2	Running	0	51m
metrics-server-f46f56d7b-t9zqh	2/2	Running	0	51m

```
PS C:\Users\HP>
```

3) Similarly run commands for other namespaces

4) Now for creating namespace

kubectl create namespace keepdoingit - create namespace

```
PS C:\Users\HP> kubectl create namespace keepdoingit
namespace/keepdoingit created
PS C:\Users\HP>
```

5) **kubectl get namespace**

```
PS C:\Users\HP> kubectl get namespace
```

NAME	STATUS	AGE
default	Active	64m
keepdoingit	Active	2m9s
kube-node-lease	Active	64m
kube-public	Active	64m
kube-system	Active	64m

```
PS C:\Users\HP>
```

6) Now run pod in this namespace – so run command helps us to create pod

kubectl run nginx --image=nginx -n keepdoingit

```
PS C:\Users\HP> kubectl run nginx --image=nginx -n keepdoingit
pod/nginx created
PS C:\Users\HP>
```

The command **'kubectl run nginx --image=nginx -n keepdoingit'** creates a new Kubernetes pod named **'nginx'** in the namespace **'keepdoingit'**, using the **'nginx'** image.

Here's a breakdown:

- **'kubectl run'**: Command to create a new pod.

- `nginx`: Name of the pod.

- `--image=nginx`: Specifies the container image to use (in this case, the official Nginx image).

- `-n keepdoingit`: Indicates the namespace where the pod will be created, which is `keepdoingit`.

This command essentially sets up an **Nginx web server** in the specified namespace.

kubectl run dhoni --image=nginx -n keepdoingit

```
PS C:\Users\HP> kubectl run dhoni --image=nginx -n keepdoingit
pod/dhoni created
PS C:\Users\HP>
```

7) **kubectl get pods -n keepdoingit** = displays a list of all pods in the Kubernetes namespace "keepdoingit"

```
PS C:\Users\HP> kubectl get pods -n keepdoingit
NAME      READY   STATUS    RESTARTS   AGE
dhoni     1/1     Running   0           107s
nginx     1/1     Running   0           7m24s
PS C:\Users\HP>
```

8) **kubectl delete pod dhoni -n keepdoingit** – deleted pod dhoni

```
PS C:\Users\HP> kubectl delete pod dhoni -n keepdoingit
pod "dhoni" deleted
PS C:\Users\HP>
```

9) **kubectl get pods -n keepdoingit** - check running pods in namespace keepdoingit

```
PS C:\Users\HP> kubectl get pods -n keepdoingit
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           15m
PS C:\Users\HP>
```

10) Now we can see after deleting this dhoni pod, controller should auto create the new pod but it has not so for this **pain** we will see how controller will create automatically new pods after pods die

