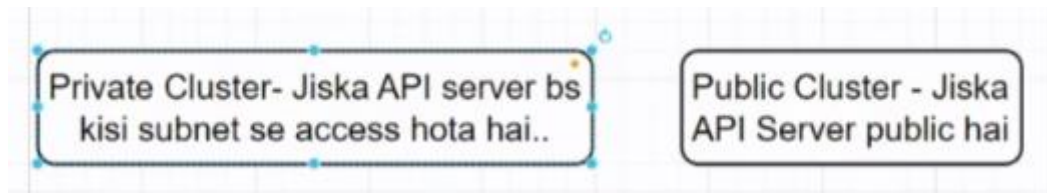


26 Oct 2024

NOTE : 1) Difference between private and public cluster.

2) Difference between kubectl create and kubectl apply command even they almost do the same work = create command creates the new resource as one time creation but apply command creates and updates the already created resources.



1) Create folder “3) 26 October Kubernetes” and copy data of previous folder

AGENDA – Network policy, PV, PVC

1)

The screenshot shows the Visual Studio Code interface. The Explorer panel on the left shows a folder named '3) 26 OCTOBER KUBERNETES' containing a subfolder 'k8scluster' with files like '.terraform', '.terraform.lock.hcl', 'main.tf', 'namespace.txt', 'namespace.yaml', 'networkpolicy.yaml', 'nginxpod.yaml', 'np.txt', 'pod.txt', 'providers.tf', 'terraform.tfstate', and 'terraform.tfstate.backup'. The main editor shows a file named 'firefoxpod.yaml' with the following content:

```
k8scluster > ! firefoxpod.yaml
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: firefox-pod
5    namespace: default
6  spec:
7    containers:
8    - image: linuxserver/firefox
9      name: fire-container
10     ports:
11     - containerPort: 3000
```

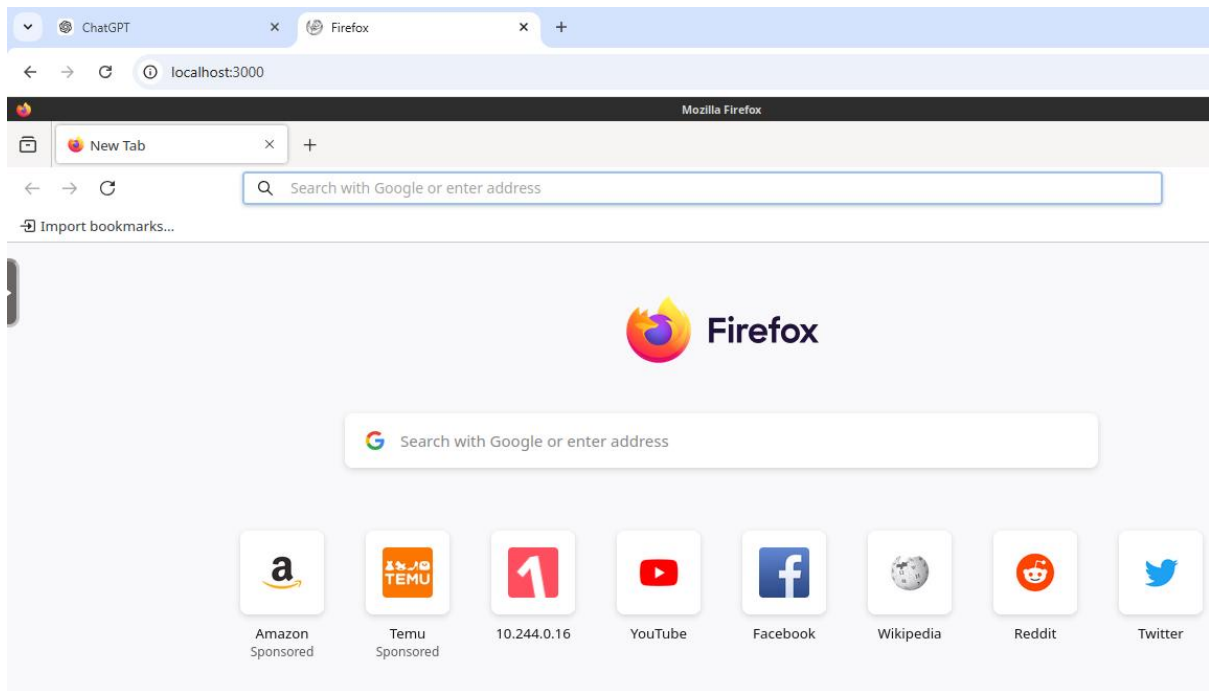
The TERMINAL panel at the bottom shows the command 'kubectl get pods' and its output:

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
firefox-pod         1/1     Running   0           5h14m
nginx-pod-with-label 1/1     Running   0           3h
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster>
```

2) **Kubectl get pods**

3) **kubectl port-forward firefox-pod 3000:3000** = It created a new tunnel from our laptop to that firefox pod

4) localhost:3000



5) Open new terminal in vscode

kubectl apply -f networkpolicy.yaml

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl apply -f networkpolicy.yaml
Warning: resource networkpolicies/nginx-network-policy is missing the kubect1.kubernetes.io/last-applied-configuration annotation which is
required by kubectl apply. kubectl apply should only be used on resources created declaratively by either kubectl create --save-config or
kubectl apply. The missing annotation will be patched automatically.
networkpolicy.networking.k8s.io/nginx-network-policy configured
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster>
```

6) **kubectl get networkpolicy**

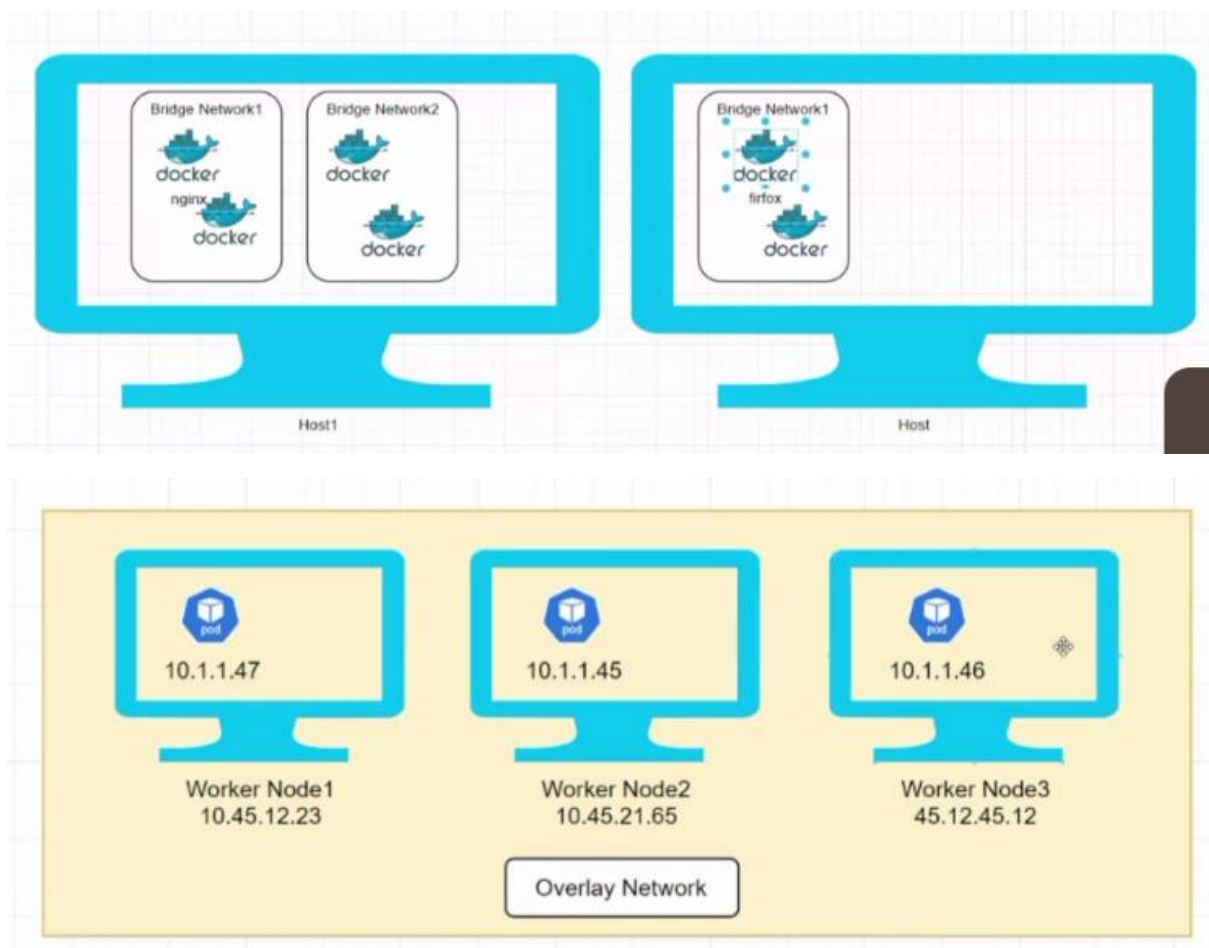
```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl get networkpolicy
NAME                POD-SELECTOR  AGE
nginx-network-policy  papa=dhondhu  3h56m
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster>
```

NOTE : 1) One pod in one worker node can communicate to another pod in another worker node by their ips.

2) The responsibility of giving ips in cluster to pods is of network plugin



3) Overlay network in docker



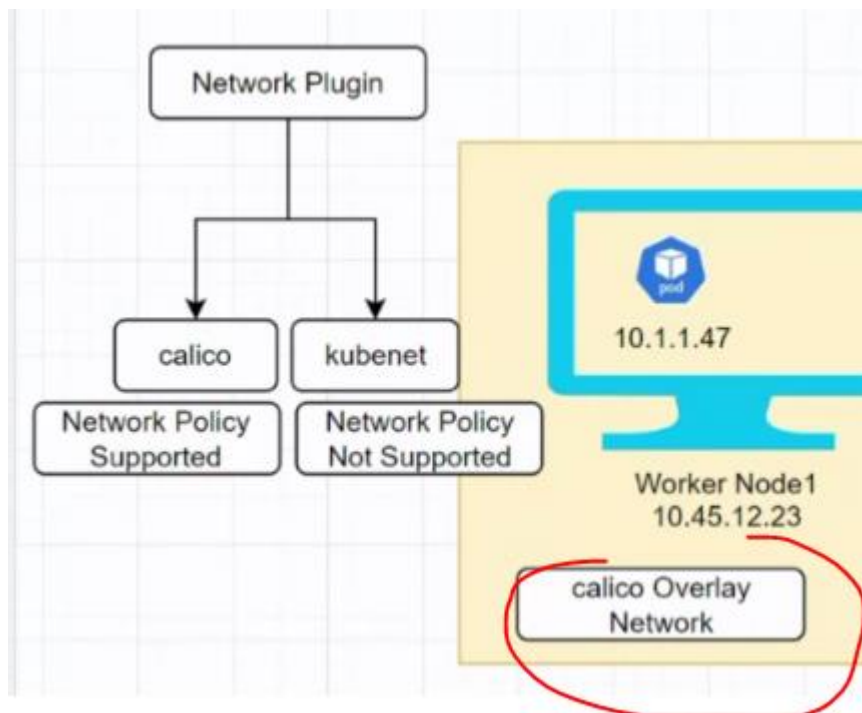
Docker में **overlay network** एक प्रकार का नेटवर्क होता है जो कंटेनरों को अलग-अलग होस्ट्स (machines) पर एक दूसरे से जोड़ता है, जैसे वे एक ही होस्ट पर हों।

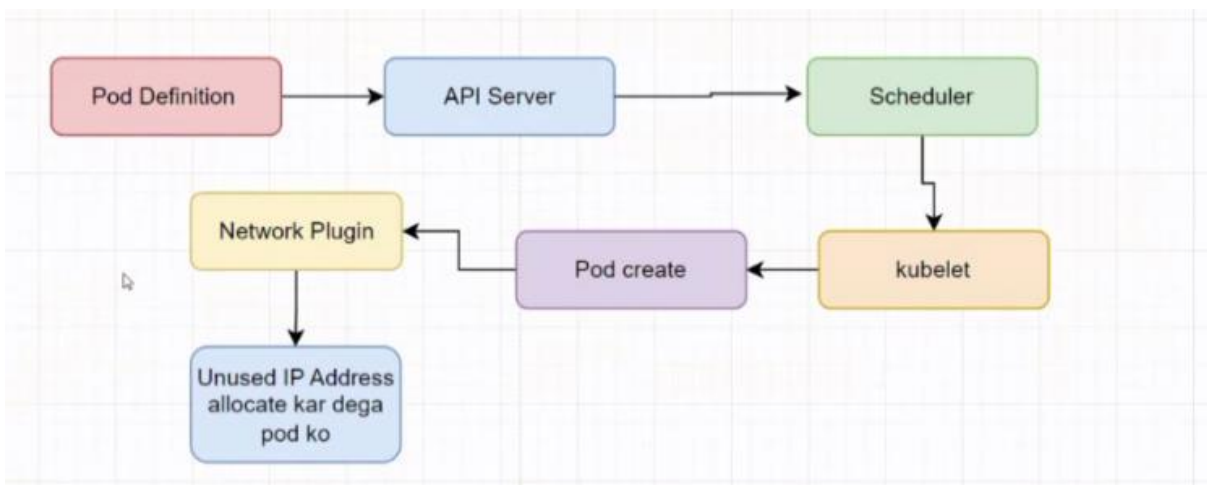
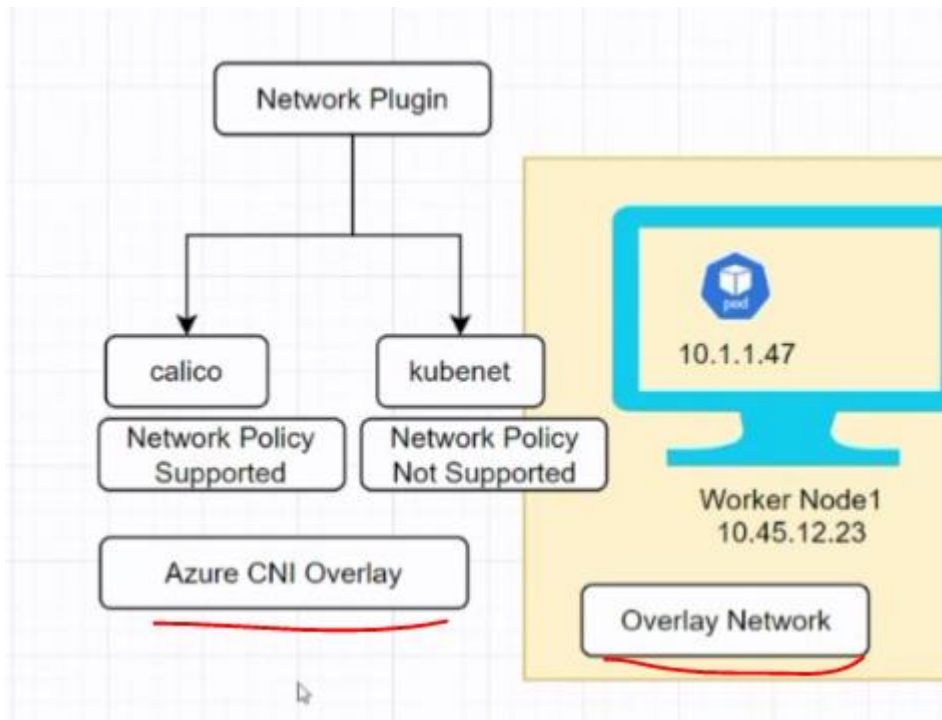
जब हम Docker Swarm या Kubernetes जैसी orchestration tools का उपयोग करते हैं, तो overlay network कंटेनरों को एक दूसरे से securely और efficiently communicate करने की सुविधा देता है, चाहे वे किसी भी physical machine पर हों।

साधारण शब्दों में कहें तो, यह एक virtual नेटवर्क होता है जो विभिन्न मशीनों पर running containers के बीच कनेक्शन बनाता है।

****Example:**** अगर आपके पास दो Docker host हैं और दोनों पर containers चल रहे हैं, तो overlay network इन कंटेनरों को आपस में संवाद करने का तरीका प्रदान करता है, जैसे वे एक ही host पर चल रहे हों।

7)





NOTE : 1) kubenet network do not supports network policy. Donot allow site to site vpn connectivity

2) AZURE cni with calico allows to set network policy & allow site to site vpn connectivity

3) Difference between kubenet and azure cni (azure container networking interface)

+++++

AGENDA- Create k8s cluster

1)

Home > Kubernetes services >

Create Kubernetes cluster

Basics Node pools Networking Integrations Monitoring Security Advanced Tags Review + create

Project details

Select a subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * ⓘ Free Trial

Resource group * ⓘ (New) rgrocket
[Create new](#)

Cluster details

Cluster preset configuration * Dev/Test

To quickly customize your Kubernetes cluster, choose one of the preset configurations above. You can modify these configurations at any time.
[Compare presets](#)

Kubernetes cluster name * ⓘ k8srocket

Region * ⓘ (Europe) Poland Central

Create Kubernetes cluster

Basics **Node pools** Networking Integrations Monitoring Security Advanced Tags Review + create

Node pools

In addition to the required primary node pool configured on the Basics tab, you can also add optional node pools to handle a variety of workloads [Learn more](#)

+ Add node pool Delete

| <input checked="" type="checkbox"/> | Name | Mode | Node size | OS SKU | Node count | Availat |
|-------------------------------------|-----------|--------|-----------------------|--------|------------|---------|
| <input checked="" type="checkbox"/> | agentpool | System | Standard_DS2_v2 (...) | Ubuntu | 1 | None |

Enable virtual nodes

Home > Kubernetes services >

Create Kubernetes cluster

Basics Node pools **Networking** Integrations Monitoring Security Advanced Tags Review + create

Container networking

Network configuration ⓘ

- ☒ **Azure CNI Overlay**
Assigns pod IP addresses from a private IP space. Best for scalability
- ☐ **Azure CNI Node Subnet**
Previously named Azure CNI. Assigns pod IP addresses from your host VNet. Best for workloads where pods must be reachable by other VNet resources
- ☐ **kubenet**
Older, route table-based Overlay with limited scalability. Not recommended for most clusters

Bring your own Azure virtual network ⓘ ☐

DNS name prefix * ⓘ

Enable Cilium dataplane and network policy ⓘ ☐

Network policy * ⓘ

- ☐ **None**
Allow all ingress and egress traffic to the pods
- ☒ **Calico**
Open-source networking solution. Best for large-scale deployments with strict security requirements
- ☐ **Azure**

Kubernetes services

| | |
|--------------------|--|
| Encryption type | Encryption at-rest with a platform-managed key |
| Virtual node pools | Not enabled |

Node pools

| | |
|---------------------|-----------------|
| Node pools | 1 node pool |
| Kubernetes versions | 1.30.6 |
| Node sizes | Standard_DS2_v2 |

Configuration

| | |
|----------------------------------|-------------------------------------|
| Kubernetes version | 1.30.6 |
| Auto Upgrade Type | Patch |
| Automatic upgrade scheduler | Every week on Sunday (recommended) |
| Node security channel type | Node Image |
| Security channel scheduler | Every week on Sunday (recommended) |
| Authentication and Authorization | Local accounts with Kubernetes RBAC |
| Local accounts | Enabled |

Extensions + applications

No extensions installed

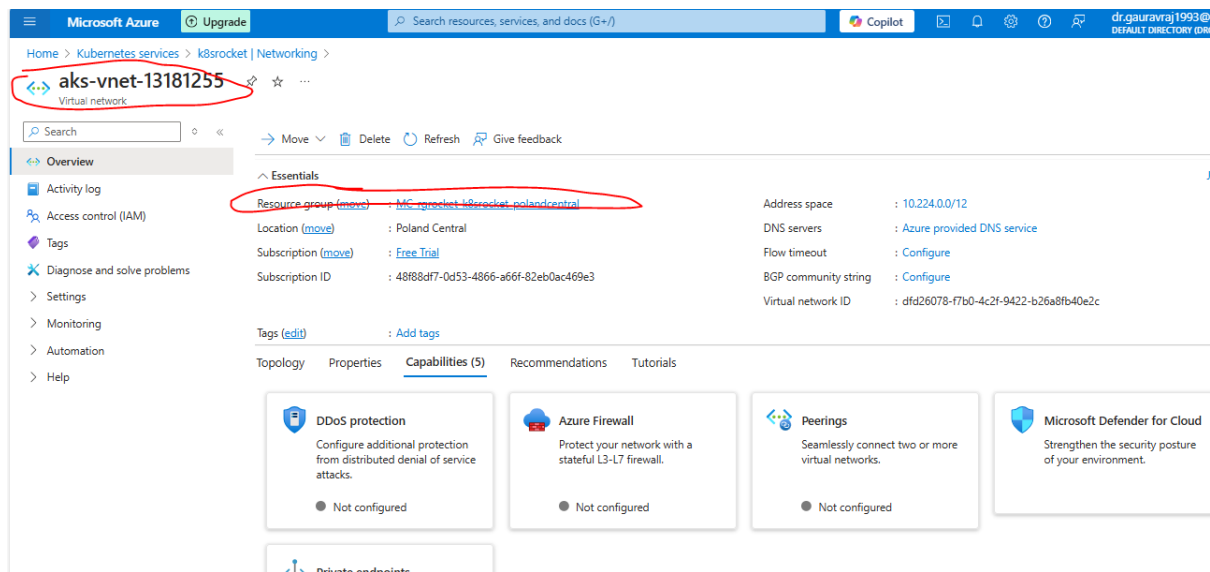
Networking

| | |
|--|--|
| API server address | k8srocket-dns-yy5j197x.hcp.polandcentral.azmk8s.io |
| Network configuration | Azure CNI Overlay |
| Pod CIDR | 10.244.0.0/16 |
| Service CIDR | 10.0.0.0/16 |
| DNS service IP | 10.0.0.10 |
| Cilium dataplane | Not enabled |
| Network Policy | Calico |
| Load balancer | Standard |
| Private cluster | Not enabled |
| Authorized IP ranges | Not enabled |
| Application Gateway ingress controller | Not enabled |

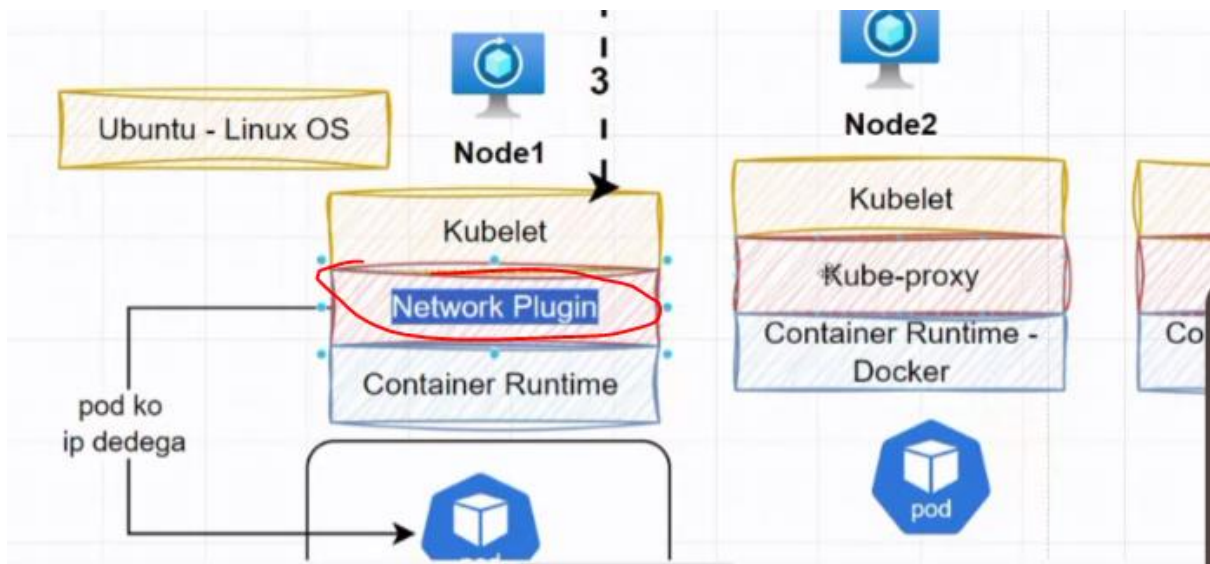
Integrations

| | |
|-----------------------|-------------|
| Container insights | Not enabled |
| Workspace resource ID | - |
| Service Mesh - Istio | Not enabled |

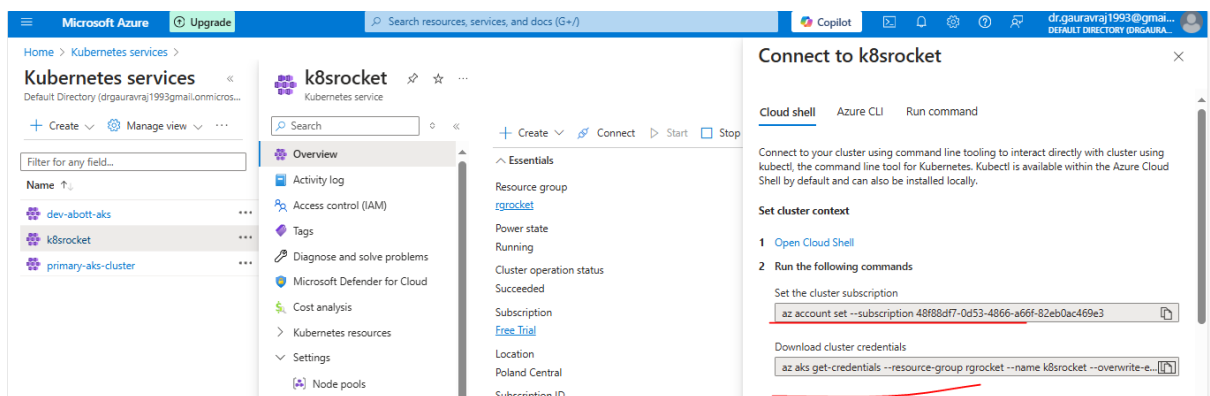
2) So above we can see calico network policy which supports network policy and Along with aks cluster a vnet is also created by azure itself in calico.



3) Now in diagram mention cni plugin or network plugin instead of kubeproxy which assigns ip to pod



4) Now since cluster is created so connecting it on our computer




```
PS C:\4) KUBERNETES\3) 26 October Kubernetes> cd .\k8scluster\
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> az account set --subscription 48f88df7-0d53-4866-a66f-82eb0ac469e3
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> az aks get-credentials --resource-group rgrocket --name k8srocket --overwrite-existing
Merged "k8srocket" as current context in C:\Users\HP\.kube\config
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> |
```

5) **kubectl apply -f firefoxpod.yaml** = create firefox pod

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl apply -f firefoxpod.yaml
pod/firefox-pod created
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> |
```

6) **kubectl apply -f nginxpod.yaml** = create nginx pod

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl apply -f nginxpod.yaml
pod/nginx-pod-with-label created
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> |
```

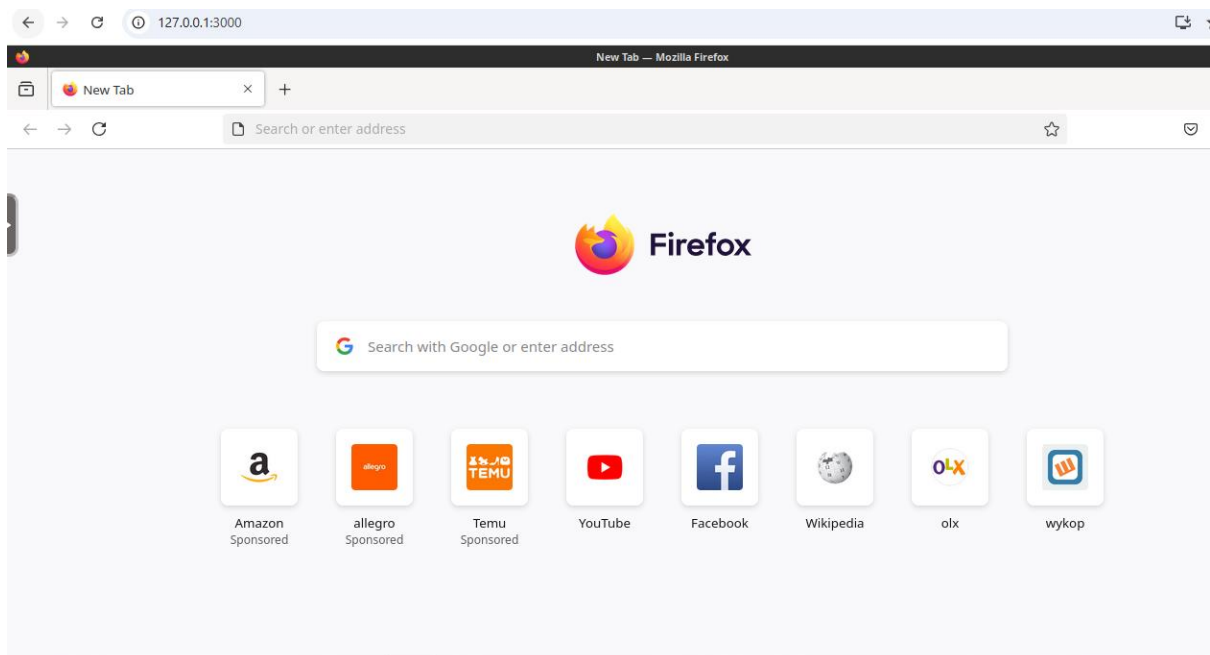
7) **kubectl get pods**

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
firefox-pod         1/1     Running   0           119s
nginx-pod-with-label 1/1     Running   0           33s
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> |
```

8) **kubectl port-forward firefox-pod 3000:3000** = port forwarding firefox from our local till pod or creating tunnel

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl port-forward firefox-pod 3000:3000
Forwarding from 127.0.0.1:3000 -> 3000
Forwarding from [::1]:3000 -> 3000
|
```

9) **127.0.0.1:3000** = Run in browser



10) Change to another terminal in vscode

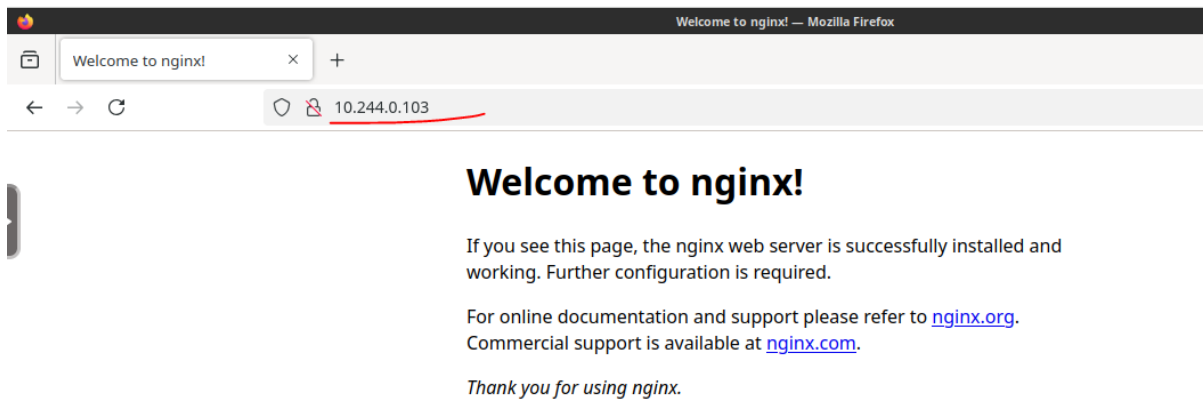
kubectl get pods -o wide

```
PS C:\4\ KUBERNETES\3) 26 October Kubernetes> kubectl get pods -o wide
```

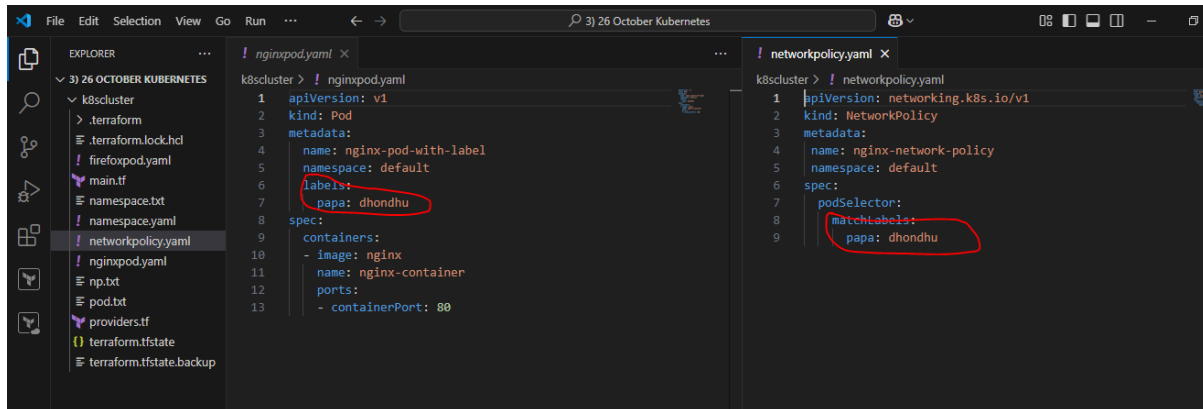
| NAME | READY | STATUS | RESTARTS | AGE | IP | NODE | NOMINATED NODE | READINESS |
|----------------------|-------|---------|----------|-------|--------------|-----------------------------------|----------------|-----------|
| GATES | | | | | | | | |
| firefox-pod | 1/1 | Running | 0 | 8m33s | 10.244.0.40 | aks-agentpool-42392433-vmss000000 | <none> | <none> |
| nginx-pod-with-label | 1/1 | Running | 0 | 7m7s | 10.244.0.103 | aks-agentpool-42392433-vmss000000 | <none> | <none> |

```
PS C:\4\ KUBERNETES\3) 26 October Kubernetes>
```

11) 10.244.0.103= run nginx pod ip in firefox



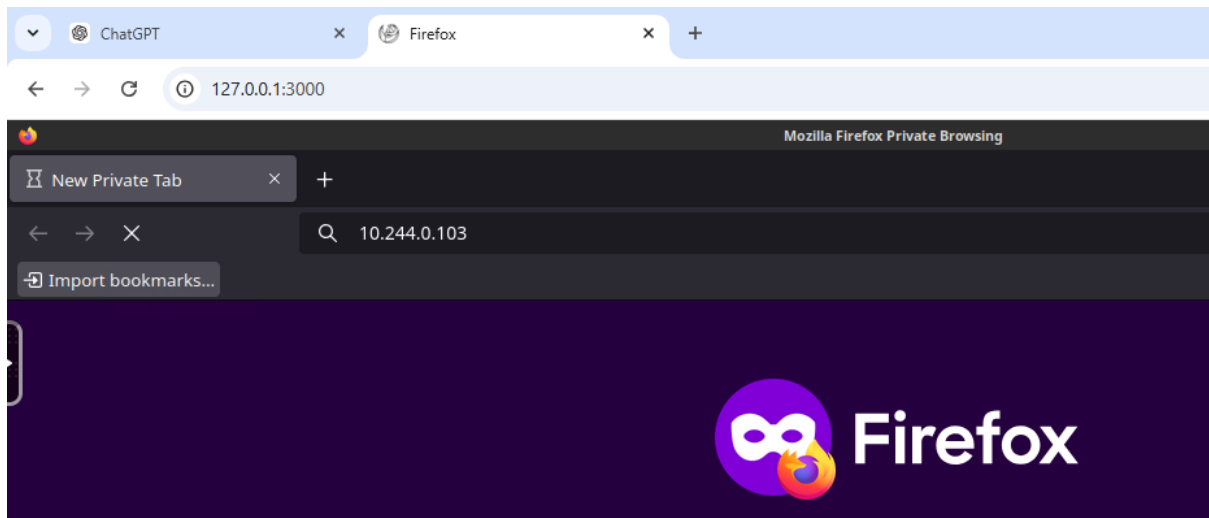
12) Now applying network policy with **labels**



13) **kubectl apply -f networkpolicy.yaml** = apply network policy

```
PS C:\4\ KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl apply -f networkpolicy.yaml
networkpolicy.networking.k8s.io/nginx-network-policy created
PS C:\4\ KUBERNETES\3) 26 October Kubernetes\k8scluster>
```

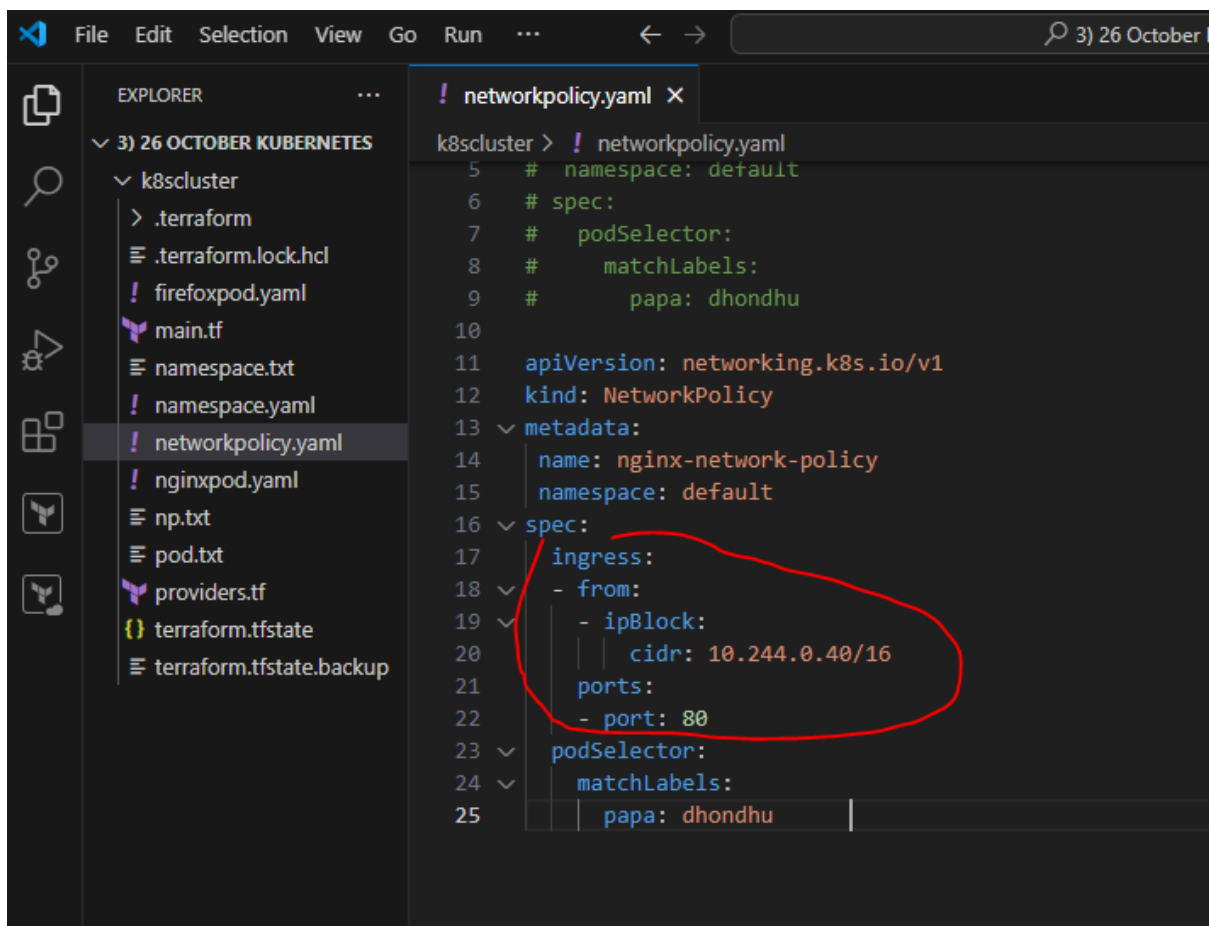
14) **10.244.0.103** = Now run ip of nginx in firefox then it will not run, which shows that network policy has been set, which means all ports on nginx pod has been closed



+++++

AGENDA – OPENING A SPECIFIC PORT ON NGINX POD

1) Now writing network policy yaml to open port 80



2) **kubectl apply -f networkpolicy.yaml**

```

nginx-pod-with-label 1/1 Running 0 15m 10.244.0.103 aks-agentpool-42392433-v
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl apply -f networkpolicy.yaml
networkpolicy.networking.k8s.io/nginx-network-policy configured
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster>

```

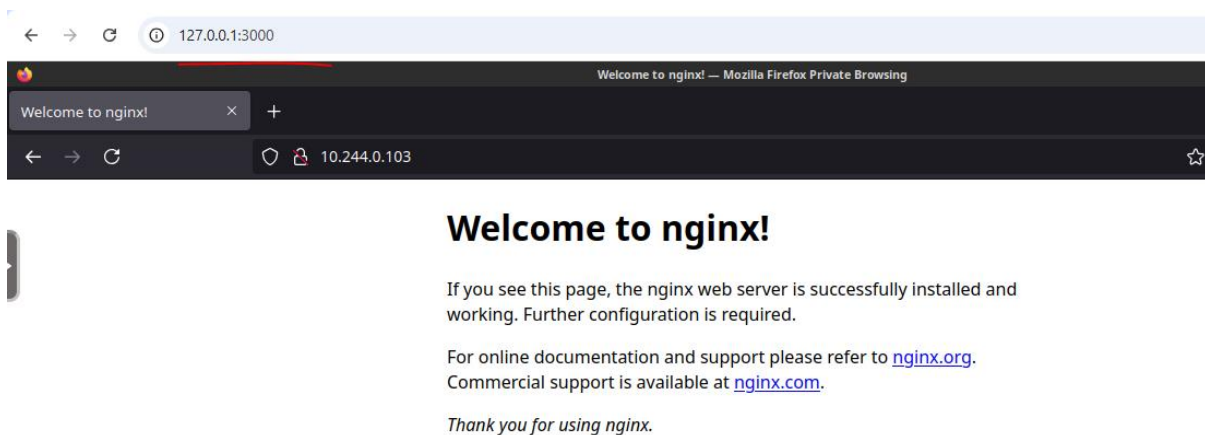
3) kubectl get networkpolicy

```

networkpolicy.networking.k8s.io/nginx-network-policy configured
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl get networkpolicy
NAME          POD-SELECTOR  AGE
nginx-network-policy  papa=dhondhu  56m
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster>

```

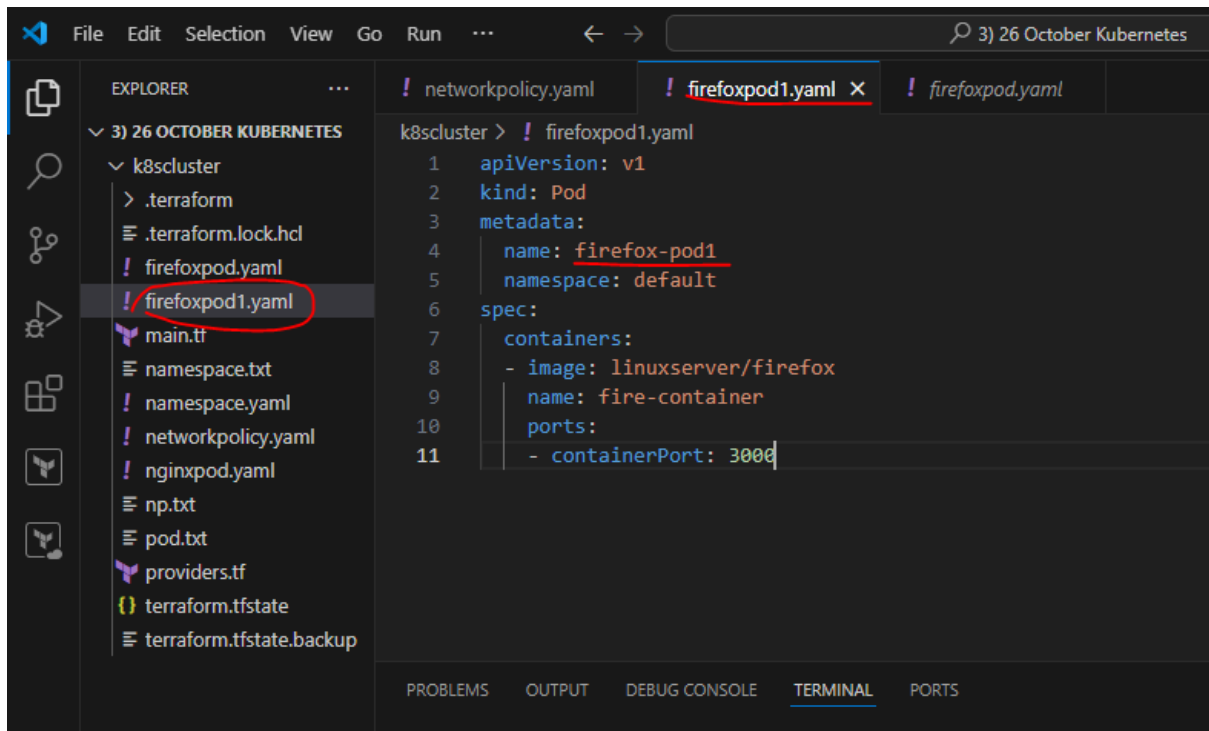
4) <http://10.244.0.103/> = Again run ip of nginx which shows that 80 port opened to run nginx or access nginx



+++++

AGENDA – CREATING FIREFOX POD AGAIN

1) Create “firefoxpod1.yaml” file and write code into it.



2) **kubectl apply -f firefoxpod1.yaml** = creating new firefox pod

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl apply -f firefoxpod1.yaml
pod/firefox-pod1 created
```

3) **kubectl get pods**

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl get pods
```

| NAME | READY | STATUS | RESTARTS | AGE |
|----------------------|-------|---------|----------|-----|
| firefox-pod | 1/1 | Running | 0 | 86m |
| firefox-pod1 | 1/1 | Running | 0 | 7s |
| nginx-pod-with-label | 1/1 | Running | 0 | 84m |

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster>
```

4) **kubectl get pods -o wide**

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl get pods -o wide
```

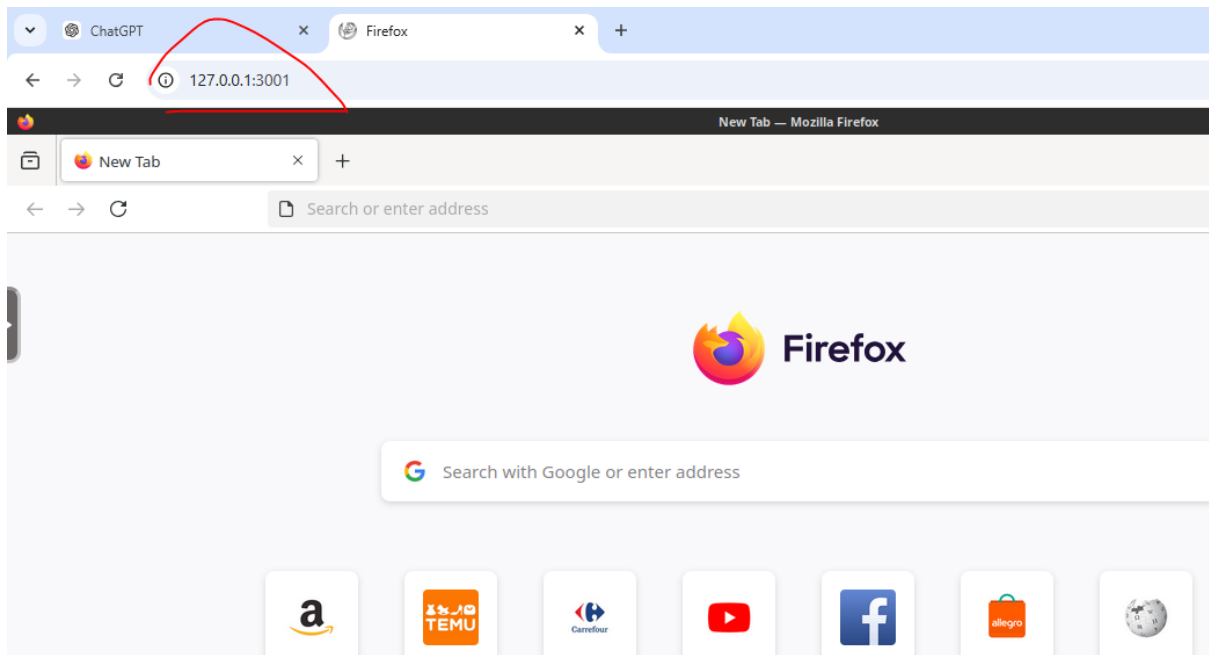
| NAME | READY | STATUS | RESTARTS | AGE | IP | NODE | NOMINATED NODE | READINESS G |
|----------------------|-------|---------|----------|------|--------------|-----------------------------------|----------------|-------------|
| firefox-pod | 1/1 | Running | 0 | 88m | 10.244.0.40 | aks-agentpool-42392433-vmss000000 | <none> | <none> |
| firefox-pod1 | 1/1 | Running | 0 | 2m5s | 10.244.0.245 | aks-agentpool-42392433-vmss000000 | <none> | <none> |
| nginx-pod-with-label | 1/1 | Running | 0 | 86m | 10.244.0.103 | aks-agentpool-42392433-vmss000000 | <none> | <none> |

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster>
```

5) **kubectl port-forward firefox-pod1 3001:3000** = creating tunnel between our laptop and firefox pod

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\k8scluster> kubectl port-forward firefox-pod1 3001:3000
Forwarding from 127.0.0.1:3001 -> 3000
Forwarding from [::1]:3001 -> 3000
```

6) **127.0.0.1:3001** = run new firefox in browser



7) **10.244.0.103** = run nginx ip in new firefox and its running because we have opened cidr of all ranges in networkpolicy.yaml file



Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

NOTE : 1) Where is network policy applied in k8s ? = On pod

2) Can we apply network policy on node? = no

3) If network policy is not working after applying, then what we will check? = Network plugin

8) Now deleting both firefox pods

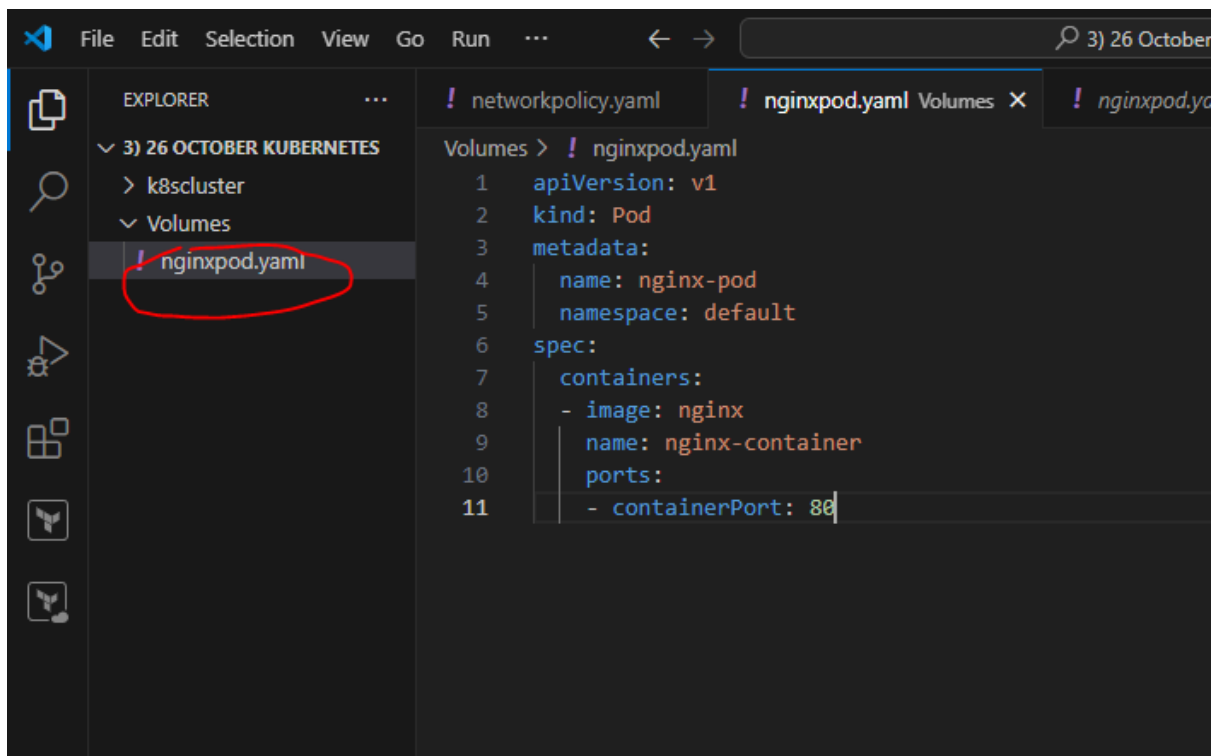
kubectl delete pod firefox-pod firefox-pod1

```
nginx-pod-with-label 2/2 Running 0 125m
PS C:\4) KUBERNETES\3) 26 October Kubernetes> kubectl delete pod firefox-pod firefox-pod1
pod "firefox-pod" deleted
pod "firefox-pod1" deleted
PS C:\4) KUBERNETES\3) 26 October Kubernetes> |
```


+++++

AGENDA – PV and PVC

1) Create folder Volumes and write nginx pod yaml file



2)



3) **kubectl apply -f nginxpod.yaml** = create pod

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes> kubectl apply -f nginxpod.yaml
pod/nginx-pod created
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes> |
```

4) **kubectl exec nginx-pod -c nginx-container -i -t -- bash** = exec command is used to enter in container inside pod

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes> kubectl exec nginx-pod -c nginx-container -i -t -- bash
root@nginx-pod:/# ls
```

cd /usr/share/nginx/html

```
root@nginx-pod:/# cd /usr/share/nginx/html
root@nginx-pod:/usr/share/nginx/html# ls
50x.html index.html
root@nginx-pod:/usr/share/nginx/html# |
```

5) **touch loveletter.txt** = create file

```
root@nginx-pod:/usr/share/nginx/html# touch loveletter.txt
root@nginx-pod:/usr/share/nginx/html# ls
50x.html index.html loveletter.txt
root@nginx-pod:/usr/share/nginx/html# |
```

NOTE : 1) 1 container = 1 process means 1 service runs in 1 container

6) **KILL 1** = 1 is like process id of container to kill it


```

root@nginx-pod:/usr/share/nginx/html# kill 1
root@nginx-pod:/usr/share/nginx/html# command terminated with exit code 137
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes>

```

7) **kubectl get pods**

```

PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes> kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
nginx-pod            1/1     Running   1 (34s ago) 72m
nginx-pod-with-label 1/1     Running   0           3h44m
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes>

```

8) So above restart shows that when we killed container in pod then pod had the responsibility to restart the container on its own, which we can see it has done it.

9) **kubectl exec nginx-pod -c nginx-container -i -t -- bash**

cd /usr/share/nginx/html

touch dhondhu.txt

```

root@nginx-pod:/usr/share/nginx/html# touch dhondhu.txt
root@nginx-pod:/usr/share/nginx/html# ls
50x.html dhondhu.txt index.html
root@nginx-pod:/usr/share/nginx/html#

```

Kill 1 = container ko maar diya ya band ho gaya

```

root@nginx-pod:/usr/share/nginx/html# kill 1
root@nginx-pod:/usr/share/nginx/html# command terminated with exit code 137
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes>

```

kubectl get pods

```

PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes> kubectl get pods
NAME                READY   STATUS    RESTARTS   AGE
nginx-pod            1/1     Running   2 (52s ago) 91m
nginx-pod-with-label 1/1     Running   0           4h3m
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes>

```

kubectl get pods -w

```

PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes> kubectl get pods -w
NAME                READY   STATUS    RESTARTS   AGE
nginx-pod            1/1     Running   2 (103s ago) 92m
nginx-pod-with-label 1/1     Running   0           4h4m

```

10) container mara to k8s ab use restart krega

11) So after restarting we can see now the container which is created by pod do not have dhondhu.txt file that we had created in previous container

kubectl exec nginx-pod -c nginx-container -i -t -- bash

cd /usr/share/nginx/html

ls

```
PS C:\4) KUBERNETES\3) 26 October Kubernetes\Volumes> kubectl exec nginx-pod -c nginx-container -i -t -- bash
root@nginx-pod:/# cd /usr/share/nginx/html
root@nginx-pod:/usr/share/nginx/html# ls
50x.html index.html
root@nginx-pod:/usr/share/nginx/html#
```

12) Now it's a great pain that file is deleted or lost in new container so basically data is lost so we will find solution for the same.